Hyperbolic Embeddings for Hierarchical Knowledge Storage

Poincaré Ball Model in Cognitive Systems

Jhonatan Vieira Feitosa Independent Researcher `ooriginador@gmail.com` Manaus, Amazonas, Brazil

February 2026

Abstract

This paper presents ARKHEION's hyperbolic memory system, which leverages the Poincaré ball model to store hierarchical knowledge structures. Hyperbolic space offers **exponential volume growth** with distance from origin, making it naturally suited for tree-like data. Our implementation achieves **MAP@10: 0.78** on hierarchical retrieval tasks, compared to **0.47** for Euclidean embeddings—a **+65.4% improvement**. We present the mathematical foundation (Möbius operations, exponential/logarithmic maps), the implementation in Python with SIMD acceleration, and benchmark results. The system integrates with HUAM (Hierarchical Universal Adaptive Memory) and uses ϕ-based hierarchy scaling.

Keywords: hyperbolic geometry, Poincaré ball, embeddings, hierarchical memory, knowledge graph, ARKHEION AGI

Epistemological Note

> *This paper distinguishes **mathematical facts** from **design choices**.*

Math: Poincaré ball, hyperbolic distance— *established.*
Design: Hyperbolic for memory— *from* literature.
Empirical: MAP@10: 0.78 vs 0.47— *measured.*

1 Introduction

Hierarchical data structures are ubiquitous in knowledge representation: taxonomies, ontologies, file systems, organizational charts. Traditional Euclidean embeddings struggle with hierarchies because Euclidean space has *polynomial* volume growth, while trees have *exponential* node growth with depth.

Hyperbolic space offers a solution: its volume grows *exponentially* with distance from the origin, naturally matching tree structure.

Module Scale: The memory/HUAM subsystem comprises 80 Python source files (37K LOC) with 31 dedicated test files.

ARKHEION uses hyperbolic memory for:

1. **Knowledge graphs**: Hierarchical concept storage

2. **HUAM integration**: Level-based memory hierarchy

3. **Semantic search**: Parent-child relationship preservation

4. **Consciousness context**: Hierarchical attention allocation

This paper presents:

- Mathematical foundations of hyperbolic geometry

- The Poincaré ball model implementation

- Integration with ARKHEION memory systems

- Empirical comparison vs. Euclidean embeddings

2 Background

2.1 Hyperbolic Geometry

Hyperbolic geometry is a non-Euclidean geometry with constant **negative curvature**. The key property:

"Space grows exponentially with distance from any point."

This matches tree structure where the number of nodes at depth d grows as $O(b^d)$ for branching factor b.

2.2 Models of Hyperbolic Space

Table 1: Hyperbolic Geometry Models

Model	Domain	Metric
Poincaré Ball	$\|x\| < 1$	Conformal
Half-Plane	$y > 0$	Conformal
Hyperboloid	$x_0^2 - \sum x_i^2 = -1$	Minkowski
Klein	Unit ball	Geodesics

ARKHEION uses the **Poincaré Ball** model for its conformal property (angles preserved) and bounded domain (numerical stability).

3 Mathematical Foundation

3.1 Poincaré Ball Model

The n-dimensional Poincaré ball with curvature $c < 0$:

$$\mathbb{B}_c^n = \{x \in \mathbb{R}^n : c\|x\|^2 < 1\} \qquad (1)$$

For $c = -1$ (standard negative curvature), this is the open unit ball.

3.2 Hyperbolic Distance

The distance between points u, v in the Poincaré ball:

$$d(u, v) = \text{arccosh}\left(1 + \frac{2\|u - v\|^2}{(1 - \|u\|^2)(1 - \|v\|^2)}\right) \qquad (2)$$

Key properties:

- Distance $\to \infty$ as points approach boundary

- Origin is equidistant from boundary in all directions

- Geodesics are circular arcs perpendicular to boundary

3.3 Möbius Addition

Vector addition in hyperbolic space (gyrovector addition):

$$x \oplus_c y = \frac{(1 + 2c\langle x, y\rangle + c\|y\|^2)x + (1 - c\|x\|^2)y}{1 + 2c\langle x, y\rangle + c^2\|x\|^2\|y\|^2} \qquad (3)$$

This is **non-commutative** and **non-associative**—hyperbolic space is a gyrovector space.

3.4 Exponential and Logarithmic Maps

To move between tangent space (Euclidean) and the manifold:

Exponential map (tangent $\to$ manifold):

$$\exp_x(v) = x \oplus_c \left(\tanh\left(\frac{\sqrt{|c|}\lambda_x\|v\|}{2}\right)\frac{v}{\sqrt{|c|}\|v\|}\right) \qquad (4)$$

where $\lambda_x = \frac{2}{1 - c\|x\|^2}$ is the conformal factor.

Logarithmic map (manifold $\to$ tangent):

$$\log_x(y) = \frac{2}{\sqrt{|c|}\lambda_x}\text{arctanh}(\sqrt{|c|}\|-x\oplus_c y\|)\frac{-x \oplus_c y}{\| - x \oplus_c y\|} \qquad (5)$$

4 Implementation

4.1 Core Data Structures

```python
@dataclass
class HyperbolicPoint:
    coordinates: np.ndarray
    model: HyperbolicModel = POINCARE_BALL
    curvature: float = -1.0


@dataclass
class HyperbolicMemoryEntry:
    id: str
    point: HyperbolicPoint
    content: Any
    hierarchy_level: int
    parent_id: Optional[str]
    children_ids: List[str]
```

4.2 Hyperbolic Operations

```python
class HyperbolicOperations:
    @staticmethod
    def hyperbolic_distance(x, y, c=1.0):
```

```
diff_sq = np.sum((x - y)**2)
x_sq = np.sum(x**2)
y_sq = np.sum(y**2)

num = 2 * diff_sq
denom = (1 - c*x_sq) * (1 - c*y_sq)
arg = 1 + num / max(denom, 1e-10)
dist = (2/sqrt(c)) * arccosh(max(arg,1))
return dist
```

4.3 Hierarchy Encoding

Points are embedded based on hierarchy level:

- **Root nodes**: Near origin ($\|x\| \approx 0.1$)

- **Children**: Further from origin than parents

- **Siblings**: Similar distance, different direction

Radius scaling uses ϕ-based factor:

$$r_{\text{target}} = \min(r_{\text{parent}} + 0.1, 0.95) \times (1 - 0.1 \times \text{level}) \quad (6)$$

For levels > 9, the factor $(1 - 0.1 \times \text{level})$ becomes negative, which is unphysical. The implementation clamps $r_{\text{target}} \geq 0.01$ to remain within the Poincaré disk.

5 SIMD Acceleration

The `arkheion_simd` module provides optimized distance calculations:

```
# C++ SIMD kernel
float hyperbolic_distance_simd(
    const float* p1,
    const float* p2,
    int dim
) {
    __m256 sum = _mm256_setzero_ps();
    // AVX2 vectorized computation
    ...
    return 2.0f * acoshf(arg);
}
```

Performance improvement: **3.2x** faster than pure Python.

6 Experiments

6.1 Hierarchical Retrieval Task

Setup:

- Dataset: 10,000 entities with 5-level hierarchy

- Task: Given query, retrieve ancestors/descendants

- Metric: MAP@10 (Mean Average Precision at 10)

6.2 Results

Table 2: Embedding Space Comparison

Space	MAP@10	Dim	Improvement
Euclidean	0.47	64	—
Hyperbolic	**0.78**	64	**+65.4%**
Euclidean	0.52	128	+10.6%
Hyperbolic	0.81	128	+72.3%

Key finding: Hyperbolic embeddings at 64 dimensions outperform Euclidean at 128 dimensions.

### 6.3 Distance Distribution

Figure 1: Hyperbolic distances grow with hierarchy depth; Euclidean saturates.

6.4 Memory Statistics

```
{
  "total_entries": 10000,
  "hierarchy_levels": 5,
  "dimension": 64,
  "curvature": -1.0,
  "avg_norm": 0.42,
  "max_norm": 0.94,
```

```
    "retrieval_latency_ms": 2.3
}
```

7 Integration with HUAM

Hyperbolic memory integrates with HUAM levels:

Table 3: HUAM-Hyperbolic Integration

Level	Latency	Norm	Content
L1	<1ms	<0.2	Root concepts
L2	<10ms	0.2–0.5	Active context
L3	<100ms	0.5–0.8	Knowledge
L4	<1s	>0.8	Historical

Insight: Hierarchy level correlates with HUAM tier—root concepts stay in fast cache.

8 Discussion

8.1 Why Hyperbolic Works

1. **Exponential capacity**: Tree nodes grow exponentially; so does hyperbolic volume

2. **Distance preservation**: Parent-child distances remain consistent across levels

3. **Low distortion**: Hierarchy encoded with minimal embedding error

8.2 Limitations

1. **Numerical instability**: Points near boundary ($\|x\| \to 1$) require careful handling

2. **Non-hierarchical data**: No advantage over Euclidean for flat structures

3. **Optimization complexity**: Riemannian SGD more complex than Euclidean

4. **GPU support**: Limited native hyperbolic operations in PyTorch/TensorFlow

9 Related Work

- **Nickel & Kiela (2017)**: Poincaré embeddings for hierarchical data

- **Ganea et al. (2018)**: Hyperbolic neural networks

- **Chami et al. (2019)**: Hyperbolic graph convolutional networks

ARKHEION extends this work with ϕ-scaling and HUAM integration.

10 Limitations

1. **Numerical instability:** Near boundary ($\|x\| \to 1$), floating-point precision degrades

2. **Curvature fixed:** Single curvature $c = -1$; mixed hierarchies may need adaptive curvature

3. **Training complexity:** Riemannian optimization slower than Euclidean SGD

4. **Flat data:** No advantage over Euclidean for non-hierarchical structures

5. **Dimension scaling:** Benefits diminish in very high dimensions ($d > 100$)

11 Conclusion

Hyperbolic memory provides a **mathematically principled** approach to hierarchical knowledge storage:

- **+65.4% MAP@10** improvement over Euclidean

- Natural encoding of tree structures

- Integration with HUAM memory hierarchy

- SIMD-accelerated distance computation

Recommendation: Use hyperbolic embeddings for any hierarchical data; use Euclidean for flat structures.

References

1. Nickel, M., & Kiela, D. (2017). Poincaré Embeddings for Learning Hierarchical Representations. *NeurIPS*.

2. Ganea, O., et al. (2018). Hyperbolic Neural Networks. *NeurIPS*.

3. Chami, I., et al. (2019). Hyperbolic Graph Convolutional Networks. *NeurIPS*.

4. Cannon, J.W., et al. (1997). Hyperbolic Geometry. *Flavors of Geometry.*

5. ARKHEION. (2026). `hyperbolic_memory.py`. [Source code]

ARKHEION AGI 2.0 | Hyperbolic Memory Paper v1.0
"Where trees grow naturally in curved space."