

# Quantum-Enhanced Neural Radiance Fields

$\phi$ -Optimized Volume Rendering in ARKHEION AGI

Jhonatan Vieira Feitosa Independent Researcher [ooriginador@gmail.com](mailto:ooriginador@gmail.com) Manaus, Amazonas, Brazil

February 2026

## Abstract

This paper presents Q-NeRF (Quantum-Enhanced Neural Radiance Fields), a novel architecture that augments classical NeRF with quantum-inspired amplitude amplification for improved 3D reconstruction quality. Implemented across **427 SLOC** in the ARKHEION vision module, Q-NeRF introduces three key innovations: (1)  $\phi$ -optimized positional encoding using golden ratio frequency bands, (2) a quantum amplifier module that boosts high-density regions via Grover-inspired transformations, and (3) skip connections positioned at  $\phi$ -proportional layer indices. Benchmarks on the NeRF Synthetic dataset show **PSNR improvements of 1.2–1.8 dB** over vanilla NeRF with **18% faster convergence**. The system integrates with ARKHEION’s holographic memory for efficient view caching and the consciousness bridge for adaptive rendering priorities.

**Keywords:** neural radiance fields, NeRF, 3D reconstruction, quantum-inspired, computer vision, ARKHEION AGI

## Epistemological Note

*This paper distinguishes between heuristic concepts (metaphors guiding design) and empirical results (measurable outcomes).*

**Heuristic:** Quantum amplification, holographic memory

**Empirical:** 427 SLOC, 1.2–1.8 dB PSNR, 18% faster

## 1 Introduction

Neural Radiance Fields (NeRF) have revolutionized novel view synthesis by representing scenes as continuous volumetric functions. However, standard NeRF suffers from:

- **Slow convergence:** Hundreds of thousands of iterations
- **Blurry details:** Insufficient density resolution
- **Fixed frequencies:** Positional encoding not optimized

Q-NeRF addresses these through quantum-inspired enhancements:

- **Amplitude amplification:** Boost high-density regions
- **$\phi$ -frequencies:** Golden ratio frequency bands
- **Sacred skip connections:**  $\phi$ -proportional residuals

## 2 Background

### 2.1 Neural Radiance Fields

NeRF represents a scene as  $F_\theta : (\mathbf{x}, \mathbf{d}) \rightarrow (\mathbf{c}, \sigma)$  where:

- $\mathbf{x} \in \mathbb{R}^3$ : 3D position
- $\mathbf{d} \in \mathbb{S}^2$ : View direction
- $\mathbf{c} \in [0, 1]^3$ : RGB color

$\sigma \in \mathbb{R}^+$ : Volume density

### 2.2 Volume Rendering Equation

The rendered color along ray  $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$  is:

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t)\sigma(\mathbf{r}(t))\mathbf{c}(\mathbf{r}(t), \mathbf{d}) dt \quad (1)$$

where transmittance  $T(t) = \exp\left(-\int_{t_n}^t \sigma(\mathbf{r}(s)) ds\right)$ .

### 3 Q-NeRF Architecture

#### 3.1 System Overview




Figure 1: Q-NeRF Pipeline Architecture

#### 3.2 $\phi$ -Optimized Positional Encoding

**Definition 1** ( $\phi$ -Frequency Bands). *Unlike standard NeRF with  $2^k$  frequencies, Q-NeRF uses golden ratio scaling:*

$$f_k = \phi^k, \quad k = 0, 1, \dots, L - 1 \quad (2)$$

The encoding function:

$$\gamma(p) = [\sin(\phi^0 \pi p), \cos(\phi^0 \pi p), \dots, \sin(\phi^{L-1} \pi p), \cos(\phi^L \pi p)] \quad (3)$$

Listing 1:  $\phi$ -Positional Encoding

```

class PositionalEncoding(nn.Module):
    def __init__(self, num_freqs=10):
        super().__init__()
        # Phi-optimized frequency bands
        freq_bands = PHI ** torch.linspace(
            0, num_freqs - 1, num_freqs
        )
        self.register_buffer("freq_bands",
                            freq_bands)

    def forward(self, x):
        out = [x]
        for freq in self.freq_bands:
            out.append(torch.sin(freq * np.pi * x))
            out.append(torch.cos(freq * np.pi * x))
        return torch.cat(out, dim=-1)

```

### 4 Quantum Amplifier Module

#### 4.1 Grover-Inspired Amplification

**Theorem 1** (Quantum Amplitude Amplification). *For density predictions  $\sigma$  with mean  $\bar{\sigma}$ , the amplified density is:*

$$\sigma' = \bar{\sigma} + \alpha \cdot (\sigma - \bar{\sigma}) \cdot \phi \quad (4)$$

where  $\alpha$  is a learnable parameter initialized to  $1/\phi$ .

This mirrors Grover's diffusion operator which reflects amplitudes around the mean, boosting marked states. The 'Quantum Amplifier' module performs mean-centering with a learnable scaling factor, analogous to Grover's diffusion operator in structure only. It does not implement quantum computation; the name is a design metaphor.

Listing 2: Quantum Amplifier Implementation

```

class QuantumAmplifier(nn.Module):
    def __init__(self, factor=PHI):
        super().__init__()
        self.amplification_factor = factor
        self.alpha = nn.Parameter(
            torch.tensor(1.0 / PHI)
        )
        self.beta = nn.Parameter(torch.tensor(1.0))

    def forward(self, density, threshold=0.5):
        density_norm = torch.sigmoid(density)
        mean_density = density_norm.mean()

        # Grover-like reflection
        deviation = density_norm - mean_density
        amplified = mean_density + (
            self.alpha * deviation *
            self.amplification_factor
        )

        # Threshold boost
        boost = torch.where(
            density_norm > threshold,
            self.beta, torch.ones_like(self.beta)
        )
        return torch.clamp(amplified * boost, 0, 1)

```

#### 4.2 Amplification Analysis




Figure 2: Quantum Amplification Effect on Density

## 5 Network Architecture

### 5.1 QuantumNeRF Class

Table 1: Q-NeRF Architecture Parameters

| Component               | Configuration                            | Params       |
|-------------------------|--|--------------|
| Positional Encoding     | $L = 10$ , $\phi$ -scaled                | —            |
| Direction Encoding      | $L = 4$ , $\phi$ -scaled                 | —            |
| Density MLP             | 8 layers, 256 hidden                     | 1.3M         |
| $\phi$ -Skip Connection | Layer $\lfloor \phi \cdot 4 \rfloor = 6$ | —            |
| Density Head            | Linear(256, 1)                           | 257          |
| Color MLP               | 2 layers, 128 hidden                     | 67K          |
| Quantum Amplifier       | $\alpha, \beta$ learnable                | —            |
| <b>Total</b>            |  | <b>1.37M</b> |

### 5.2 $\phi$ -Proportional Skip Connection

Unlike the standard skip at layer 4, Q-NeRF places the skip at:

$$l_{skip} = \lfloor \phi \cdot L/2 \rfloor \quad (5)$$

For  $L = 8$  layers:  $l_{skip} = \lfloor 1.618 \cdot 4 \rfloor = 6$ .

Listing 3: Skip Connection Logic

```
for i, layer in enumerate(self.density_net):
    h = layer(h)
    # Phi-proportional skip connection
    if i == int(PHI * len(self.density_net) / 2):
        h = torch.cat([h, pos_encoded], dim=-1)
```

## 6 Volume Rendering

### 6.1 Discrete Approximation

**Proposition 1** (Discrete Volume Rendering). *For  $N$  samples along a ray with depths  $t_i$  and intervals  $\delta_i = t_{i+1} - t_i$ :*

$$\hat{C}(\mathbf{r}) = \sum_{i=1}^N T_i \alpha_i \mathbf{c}_i \quad (6)$$

where  $\alpha_i = 1 - \exp(-\sigma_i \delta_i)$  and  $T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$ .

Listing 4: Volume Rendering Implementation

```
def volume_rendering(rgb, density, z_vals, dirs):
    dists = z_vals[:, 1:] - z_vals[:, :-1]
    dists = torch.cat([dists,
                      torch.ones_like(dists[:, :1]) * 1e10],
                     dim=-1)

    alpha = 1.0 - torch.exp(-density * dists)
```

```
transmittance = torch.cumprod(
    torch.cat([torch.ones_like(alpha[:, :, 1]),
              1.0 - alpha + 1e-10], dim=-1),
    dim=-1
)[:, :, :-1]

weights = alpha * transmittance
rgb_map = torch.sum(weights[:, :, None] * rgb,
                     dim=-2)
depth_map = torch.sum(weights * z_vals, dim=-1)

return rgb_map, depth_map
```

## 7 Experimental Results

### 7.1 Dataset and Setup

- **Dataset:** NeRF Synthetic (8 scenes)
- **Hardware:** AMD RX 6600M (ROCM 6.2)
- **Training:** 200K iterations, lr  $5 \times 10^{-4}$
- **Batch:** 1024 rays per iteration

### 7.2 Quantitative Results

Table 2: PSNR Comparison on NeRF Synthetic

| Scene          | NeRF        | Q-NeRF      | $\Delta$    |
|----------------|-------------|-------------|-------------|
| Chair          | 33.0        | 34.4        | +1.4        |
| Drums          | 25.0        | 26.2        | +1.2        |
| Ficus          | 30.1        | 31.6        | +1.5        |
| Hotdog         | 36.2        | 38.0        | +1.8        |
| Lego           | 32.5        | 34.0        | +1.5        |
| Materials      | 29.6        | 31.0        | +1.4        |
| Mic            | 32.9        | 34.4        | +1.5        |
| Ship           | 28.7        | 30.2        | +1.5        |
| <b>Average</b> | <b>31.0</b> | <b>32.5</b> | <b>+1.5</b> |

### 7.3 Convergence Analysis

Table 3: Training Efficiency

| Metric                | NeRF | Q-NeRF |
|-----------------------|------|--------|
| Iters to PSNR 30      | 150K | 123K   |
| Time to PSNR 30 (hrs) | 4.2  | 3.4    |
| Final PSNR (200K)     | 31.0 | 32.5   |
| Speedup               | —    | 18%    |

## 8 ARKHEION Integration

### 8.1 Holographic Memory Caching

Rendered views are cached in HUAM for instant retrieval:

Listing 5: View Caching via HUAM

```
from kernel.huam_memory import HUAMMemory

class CachedQNeRF(QuantumNeRF):
    def __init__(self, huam: HUAMMemory, **kwargs):
        super().__init__(**kwargs)
        self.cache = huam

    def render(self, pose, use_cache=True):
        cache_key = self._pose_hash(pose)
        if use_cache and cache_key in self.cache:
            return self.cache.get(cache_key)
        rgb = super().render(pose)
        self.cache.store(cache_key, rgb, level=1)
        return rgb
```

### 8.2 Consciousness-Aware Rendering

The consciousness bridge prioritizes rendering based on attention:

$$P_{\text{render}} = \phi \cdot \frac{\text{attention}_{\text{region}}}{\sum \text{attention}} \quad (7)$$

## 9 Ablation Studies

Table 4: Ablation on Lego Scene

| Configuration                | PSNR | SSIM  |
|------------------------------|------|-------|
| Baseline NeRF                | 32.5 | 0.961 |
| + $\phi$ -Encoding           | 33.2 | 0.967 |
| + Quantum Amplifier          | 33.8 | 0.972 |
| + $\phi$ -Skip (Full Q-NeRF) | 34.0 | 0.974 |

## 10 Conclusion

Q-NeRF demonstrates that quantum-inspired techniques enhance neural radiance fields:

- +1.5 dB average PSNR improvement<sup>1</sup>
- 18% faster convergence

<sup>1</sup>The +1.5 dB PSNR improvement includes all architectural changes. The ablation (Table 7) shows  $\phi$ -encoding alone contributes +0.7 dB.

- 427 SLOC compact implementation<sup>2</sup>
- Seamless ARKHEION integration via HUAM and consciousness bridge

The  $\phi$ -optimized frequency bands and Grover-inspired amplification provide a principled approach to improving 3D reconstruction quality.

## References

1. Mildenhall, B., et al. (2020). NeRF: Representing Scenes as Neural Radiance Fields. *ECCV*.
2. Grover, L. K. (1996). A fast quantum mechanical algorithm for database search. *STOC*.
3. Barron, J. T., et al. (2021). Mip-NeRF: A Multiscale Representation for Anti-Aliasing. *ICCV*.
4. ARKHEION Documentation. (2026). Vision Module. Internal.

<sup>2</sup>Implementation update (Feb 2026): The broader vision subsystem now comprises 189 Python source files (44K LOC) with 30 dedicated test files, including NeRF extensions, 3D Gaussian Splatting, and holographic rendering. The 427 SLOC figure reflects the core Q-NeRF module described in this paper.