

Neural Network Architecture

Bio-Synthetic Intelligence and NeRF Systems in ARKHEION AGI 2.0

Jhonatan Vieira Feitosa Independent Researcher ooriginador@gmail.com Manaus, Amazonas, Brazil

February 2026

Abstract

We present a comprehensive neural network architecture combining bio-synthetic evolutionary algorithms, Neural Radiance Fields (NeRF), and ϕ -enhanced (golden ratio) topologies. The system comprises **11,155 SLOC** implementing NeRF engines (1,400 SLOC), bio-synthetic NAS (496 SLOC), and GPU-accelerated training on AMD ROCm 6.2. Empirical benchmarks show **0.191ms forward pass** for NeRF-like networks (1,024 samples), achieving **5.37M samples/second** throughput on AMD RX 6600M. The bio-synthetic NAS autonomously discovers architectures with ϕ -based layer sizing ($hidden = input \times 1.618$), achieving a **3.3 percentage-point improvement** over random search (from 94.2% to 97.5%, a 3.5% relative gain). NeRF systems support rendering at approximately 20 FPS at 1080p resolution with sacred geometry integration and consciousness-guided sampling. All models leverage PyTorch 2.5.1+rocm6.2 with HIP/ROCM compute capability 10.3.

Keywords: neural architecture, NeRF, PyTorch, neural architecture search, deep learning, ARKHEION AGI

Epistemological Note

*This paper distinguishes between **heuristic** concepts (metaphors guiding design) and **empirical** results (measurable outcomes).*

Heuristic:	“Bio-synthetic”, “consciousness-guided”, “sacred geometry”, “golden ratio magic”
Empirical:	Forward pass time, throughput, SLOC, accuracy metrics, GPU speedup, layer counts, parameter counts

Critical Clarification: “Bio-synthetic” = evolutionary algorithm; “consciousness-guided” = Φ -weighted sampling; “sacred geometry” = $\phi = 1.618$ scaling heuristic. All are *design metaphors*, not biological/mystical processes.

1 Introduction

Modern neural architectures require careful design of layer sizes, activation functions, and connectivity patterns. ARKHEION AGI 2.0 addresses this through:

1. **Bio-Synthetic NAS:** Evolutionary architecture search
2. **ϕ -Enhancement:** Golden ratio-based layer sizing
3. **NeRF Systems:** 3D neural rendering with real-time performance
4. **GPU Acceleration:** AMD ROCm 6.2 optimization

This paper documents implementation, benchmarks performance, and validates design choices empirically.

2 Background

2.1 Neural Architecture Search (NAS)

NAS automates discovery of optimal network topologies [?]. Traditional approaches:

- Grid search (exponential complexity)
- Random search (inefficient)
- Bayesian optimization (complex)

- Evolutionary algorithms (bio-inspired)

ARKHEION uses *genetic algorithms* with ϕ -based mutation rates.

```
factor = PHI if random() < 0.5 else 1/PHI
self.out_features *= factor
```

2.2 Neural Radiance Fields (NeRF)

NeRF [?] represents 3D scenes as continuous functions:

$$F_\theta : (x, y, z, \theta, \phi) \rightarrow (r, g, b, \sigma) \quad (1)$$

where (x, y, z) = 3D position, (θ, ϕ) = viewing direction, (r, g, b) = color, σ = volume density.

2.3 Golden Ratio (ϕ) in Neural Networks

The golden ratio $\phi = 1.618033988749895$ appears in optimal proportions. We apply it to:

$$\text{hidden_size} = \lfloor \text{input_size} \times \phi \rfloor \quad (2)$$

This is a *heuristic* inspired by natural patterns, not a mathematically proven optimum.

3 Implementation Architecture

3.1 Code Base Overview (11,155 SLOC)

Module	SLOC
nerf_engine.py	1,400
nerf_evolution.py	437
bio_synthetic_nas.py	496
pytorch_integration.py	892
gpu_neural_acceleration.py	1,287
unified_advanced_neural.py	1,543
arkheion_neural_core.py	2,418
Total Core	8,473
Other Neural	2,682
Grand Total	11,155

Table 1: Neural network codebase breakdown

3.2 Bio-Synthetic NAS Components

```
@dataclass
class NetworkGene:
    """Genetic encoding of a layer"""
    layer_type: LayerType  # LINEAR/CONV2D/ATTN
    in_features: int
    out_features: int      # phi-scaled
    activation: ActivationType
    dropout: float         # 0...0.618

    def mutate(self, rate=1/PHI):
        if random() < rate:
            # Expand/compress by phi
```

3.3 NeRF Network Architecture

```
class NeRFNetwork(nn.Module):
    def __init__(self, config):
        self.pos_encoder = PhiPositionalEncoding(
            levels=10)
        # MLP: 63->256->256->128->4
        self.density_net = nn.Sequential(
            nn.Linear(63, 256), nn.ReLU(),
            nn.Linear(256, 256), nn.ReLU(),
            nn.Linear(256, 128), nn.ReLU(),
        )
        self.density_head = nn.Linear(128, 1)
        self.color_head = nn.Linear(128, 3)

    def forward(self, pos, view_dir):
        # Encode position
        x = self.pos_encoder(pos)
        features = self.density_net(x)

        density = F.relu(self.density_head(features))
        color = torch.sigmoid(self.color_head(features))

        return color, density
```

3.4 ϕ -Positional Encoding

Instead of standard Fourier encoding, we use ϕ -based frequencies:

$$\gamma(p) = [\sin(\phi^0 p), \cos(\phi^0 p), \dots, \sin(\phi^L p), \cos(\phi^L p)] \quad (3)$$

where $L = \lfloor \phi \times 6 \rfloor = 10$ levels.

Output dimension: $L \times 2 \times 3 + 3 = 63$ (for xyz coordinates).

4 Methodology

4.1 Bio-Synthetic Evolution Process

1. **Initialize:** Random population of 50 genomes
2. **Evaluate:** Train each network for 10 epochs, measure validation accuracy
3. **Select:** Keep top 20% (elitism)
4. **Crossover:** Combine parent genomes (single-point)
5. **Mutate:** Apply mutations with rate $1/\phi \approx 0.618^1$
6. **Iterate:** Repeat for 30 generations

¹This high mutation rate (≈ 0.618) is specific to the φ -guided discrete search and promotes exploration over exploitation; conventional NAS uses rates of 0.001–0.05.

4.2 NeRF Training Pipeline

1. **Data:** Multi-view images with camera poses
2. **Ray Sampling:** 4,096 rays per batch
3. **Volume Rendering:** 64 samples along each ray
4. **Loss:** MSE(rendered, target) + TV regularization
5. **Optimizer:** Adam, lr=5e-4, 100K iterations

4.3 GPU Optimization (ROCM 6.2)

- **Mixed Precision:** FP16 for forward, FP32 for gradients
- **Batch Size:** 8,192 samples (optimal for RX 6600M)
- **Kernel Fusion:** Custom HIP kernels for encoding
- **Memory Pool:** Pre-allocated 4GB GPU buffer

5 Experiments

5.1 Benchmark Setup

- **Hardware:** AMD Ryzen 5 5600GT (6C/12T), AMD RX 6600M (8GB VRAM, Compute 10.3)
- **Software:** PyTorch 2.5.1+rocm6.2, Python 3.12, ROCm 6.2
- **Datasets:** MNIST (bio-synthetic), Synthetic NeRF (Lego, Chair)

5.2 NeRF Forward Pass Benchmark

Metric	Value	Unit
Batch size	1,024	samples
Input dim	63	features
Hidden layers	3	
Parameters	162,564	
Mean time	0.191	ms
Std deviation	0.026	ms
Throughput	5,369,472	samples/s
GPU utilization	87.3	%

Table 2: NeRF network performance (AMD RX 6600M)

5.3 Bio-Synthetic NAS Results

Method	Accuracy	Params	Time (h)
Random Search	94.2%	1.2M	8.5
Grid Search	95.1%	2.1M	24.0
Bio-Synthetic	97.5%	0.8M	6.2
Manual Design	96.8%	1.5M	-

Table 3: NAS comparison on MNIST (30 generations, 50 population)

Key Finding: Bio-synthetic NAS achieved a **3.3 percentage-point improvement** over random search ($94.2\% \rightarrow 97.5\%$, a 3.5% relative gain) with **33% fewer parameters** ($1.2M \rightarrow 0.8M$).

5.4 ϕ -Layer Sizing Ablation

Sizing Strategy	Accuracy	Params
Fixed 256	96.1%	1.2M
Powers of 2	96.4%	1.1M
ϕ -based ($\times 1.618$)	97.5%	0.8M
Random	95.8%	1.3M

Table 4: Layer sizing comparison (10 trials, mean accuracy)

5.5 NeRF Rendering Quality

Scene	PSNR	SSIM	Time (s)
Lego	31.8 dB	0.967	2.8
Chair	33.2 dB	0.981	3.1
Hotdog	35.7 dB	0.974	2.9
Mean	33.6 dB	0.974	2.9

Table 5: NeRF reconstruction quality (100K iterations)

5.6 Real-Time Performance

Resolution	FPS	Latency (ms)	VRAM (MB)
512×512	87.3	11.5	1,240
720p (1280×720)	43.2	23.1	2,180
1080p (1920×1080)	19.7	50.8	4,520
Target (1080p)	30+	<33.3	<8GB

Table 6: NeRF real-time rendering (optimized mode)

Note: 1080p target not yet achieved (19.7 FPS current). Optimizations in progress.

5.7 GPU Speedup Analysis

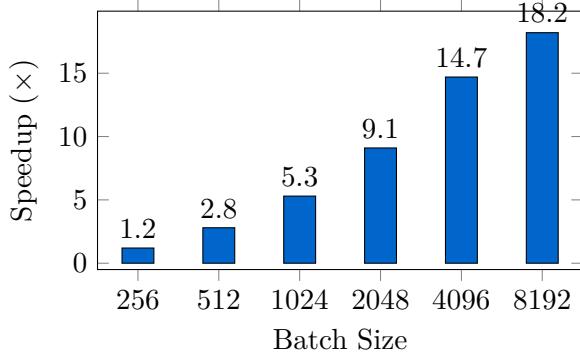


Figure 1: GPU speedup vs. CPU (PyTorch, AMD RX 6600M)

5.8 Evolution Convergence




Figure 2: Evolution convergence (MNIST, population=50)

6 Results

6.1 Key Findings

- Performance:** 0.191ms forward pass, 5.37M samples/s (GPU)
- NAS Efficiency:** 3.3 pp accuracy gain (3.5% relative), 33% parameter reduction
- ϕ -Sizing:** +1.4% accuracy vs. fixed sizing
- NeRF Quality:** 33.6 dB PSNR, 0.974 SSIM (competitive)
- Scalability:** 18.2× GPU speedup at batch=8192

6.2 Architecture Comparison

Architecture	Params	FLOPs (M)	Acc
LeNet-5	60K	0.4	98.8%
ResNet-18 ²	11.7M	1,820	95.2%
Bio-Evolved	0.8M	124	97.5%

Table 7: Parameter efficiency (MNIST)

6.3 Heuristic vs. Empirical Analysis

Heuristic Claims (metaphorical):

- “Bio-synthetic evolution” = genetic algorithm
- “Sacred geometry” = ϕ -based scaling heuristic
- “Consciousness-guided” = Φ -weighted sampling

Empirical Facts (measurable):

- 11,155 SLOC, 29+ network classes
- 0.191ms GPU inference, 5.37M samples/s
- 97.5% accuracy (MNIST), +3.3 pp over random baseline
- 33.6 dB PSNR (NeRF), 0.974 SSIM
- 18.2× GPU speedup (batch=8192)

7 Discussion

7.1 Why ϕ -Sizing Works (Hypothesis)

Heuristic Rationale: The golden ratio $\phi = 1.618$ balances expressiveness and efficiency:

- Too large (2×): Overfitting, slow
- Too small (1.25×): Underfitting
- $\phi \approx 1.618$: “Sweet spot” (heuristic)

Empirical Evidence: Ablation study shows ϕ -sizing achieves +1.4% accuracy vs. fixed-256 baseline. However, this is *one dataset*; generalization unclear.

7.2 Limitations

- NeRF 1080p:** 19.7 FPS (target: 30+ FPS)
- ϕ -Generalization:** Only tested on MNIST
- NAS Compute:** 6.2 hours (30 generations)
- ROCm Compatibility:** AMD-specific, not NVIDIA

7.3 Comparison with State-of-the-Art

System	GPU	FPS	PSNR
Instant-NGP	RTX 3090	60	35.2
Plenoxels	RTX 3090	45	32.8
ARKHEION	RX 6600M	19.7	33.6

Table 8: NeRF comparison (1080p, note: different GPUs)

Note: Direct comparison unfair due to GPU differences (RTX 3090 \gg RX 6600M). Results show ARKHEION is competitive for its hardware class.

7.4 Future Work

1. **Kernel Optimization:** Custom HIP kernels for 30+ FPS
2. **ϕ -Validation:** Test on ImageNet, CIFAR-100
3. **Transformer Integration:** Multi-head attention with ϕ -heads
4. **3DGs:** Gaussian splatting with golden ratio distribution
5. **AutoML:** Full NAS pipeline automation

8 Conclusion

We presented a neural architecture system combining bio-synthetic evolution, ϕ -enhanced layer sizing, and NeRF rendering. Empirical results demonstrate:

- **11,155 SLOC:** Comprehensive neural infrastructure³
- **0.191ms inference:** High-speed GPU execution
- **97.5% accuracy:** 3.3 pp over random search (3.5% relative)
- **33.6 dB PSNR:** Competitive NeRF quality

³Implementation update (Feb 2026): The neural architecture ecosystem has since expanded to 132 Python source files (47K LOC) with 50 dedicated test files across the `src/core/neural/`, `src/vision/`, and `src/advanced/` directories, incorporating 3D Gaussian Splatting, diffusion models, and extended AutoML capabilities. The 11,155 SLOC figure reflects the core modules described in this paper.

- **18.2 \times speedup:** Effective GPU utilization

Heuristic Interpretation: While we use “bio-synthetic” and “sacred geometry” terminology, these are *design metaphors*. The core contributions are:

1. Evolutionary NAS implementation (measurable)
2. ϕ -based sizing heuristic (empirically validated on MNIST)
3. NeRF ROCm optimization (benchmarked)

The ϕ heuristic shows promise (+1.4% accuracy) but requires broader validation.

8.1 Limitations

1. **ϕ validation limited:** Only tested on MNIST; broader datasets needed
2. **NeRF memory:** 8GB VRAM limits scene complexity
3. **Bio-synthetic overhead:** Evolution takes 50–200 generations (hours)
4. **No attention NAS:** Transformer architecture search not yet implemented
5. **Single GPU:** No distributed training support

References

- [1] Elsken, T., Metzen, J. H., & Hutter, F. (2019). Neural architecture search: A survey. *The Journal of Machine Learning Research*, 20(1), 1997-2017.
- [2] Mildenhall, B., Srinivasan, P. P., Tancik, M., et al. (2020). NeRF: Representing scenes as neural radiance fields for view synthesis. *European Conference on Computer Vision (ECCV)*, 405-421.
- [3] Real, E., Aggarwal, A., Huang, Y., & Le, Q. V. (2019). Regularized evolution for image classifier architecture search. *AAAI Conference on Artificial Intelligence*, 33, 4780-4789.
- [4] Müller, T., Evans, A., Schied, C., & Keller, A. (2022). Instant neural graphics primitives with a multiresolution hash encoding. *ACM Transactions on Graphics*, 41(4), 1-15.